Processing large-scale data efficiently: An introduction to the R package 'data.table'.

Research Data Scotland

Bayes Centre
University of Edinburgh
15th August 2024

Andreas Hoehn

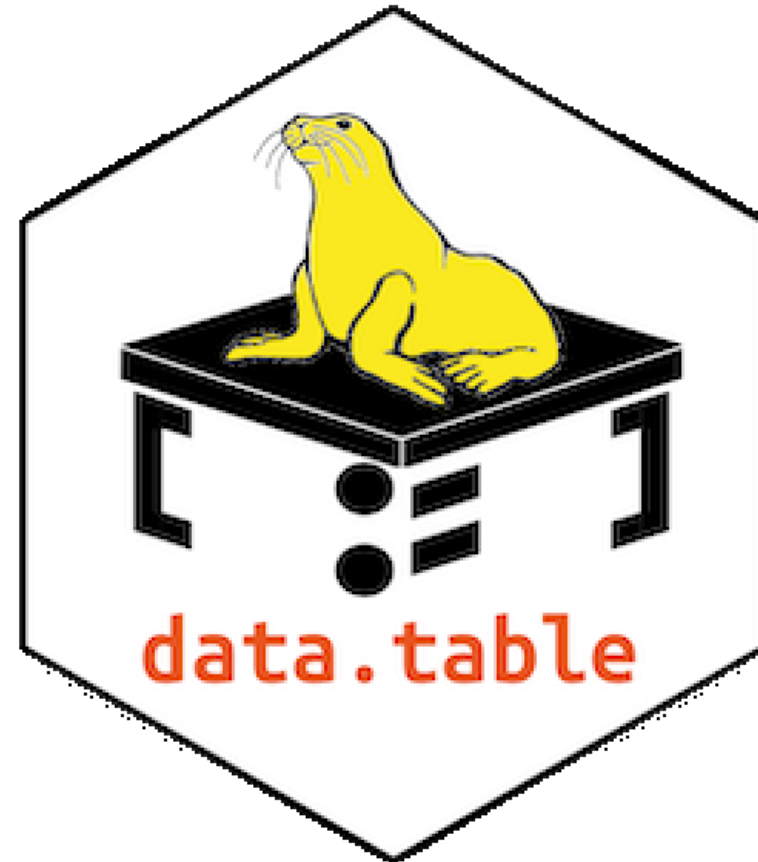University of Glasgow
andreas.hoehn@glasgow.ac.uk
AndreasXHoehn
0000-0002-7170-1205

Outline

Introducing 'data.table'

Hands-on Session

Q&A

Source: https://rdatatable.gitlab.io/data.table/

Introducing 'data.table'

What will be faster: A Ferrari or a Honda?

Introducing 'data.table'

What will be faster: A Ferrari or a Honda?

Photo by Stefano Probst on Unsplash

vs.

Photo by Brad armore on Unsplash

Introducing 'data.table'

Apple-to-apple comparisons are required!

Source: https://h2oai.github.io/db-benchmark/

Introducing 'data.table'

Scene Setting

1) Routinely collected data has always been large - and will only get larger (e.g. high-dimensional smart data, real-time data, omics data, CTGAN synthetic data)

2) Memory (RAM) and processing capacity (CPU) have remained limited resources

While R is a fantastic programming language, "standard" R (e.g. through base or tidyverse) is not an efficient/safe/futureproof/portable/updateable/….. way of working

Even worse! The limitations of "standard" R will fuel the problems that come from 1) and 2)

Introducing 'data.table'

"Standard" R's limitations

1) By default, R uses 1 core – but many more are available on most machines (→ CPU + time)

2) By default, R has a "copy-on-modify behaviour" (→ RAM + time)

3) Dependencies, portability, backward compatibility: Updates vs. old code (→ lots of time ….)

4) R Code can get long, especially when using long chains of pipes (→ even more time …)

Introducing 'data.table'

"Standard" R's limitations vs. data.table

1) By default, R uses 1 core – but many more are available on most machines ($\rightarrow$ CPU + time)
data.table parallelises whenever this is easily done, running compiled C/C++ underneath

2) By default, R has a "copy-on-modify behaviour" ($\rightarrow$ RAM + time)
data.table modifies on reference (" := ") and uses pointers, requiring less working RAM

3) Dependencies, portability, backward compatibility: Updates vs. old code ($\rightarrow$ lots of time ....)
data.table has no dependencies other than R >= Version 3.1 (10+ years old)

4) R Code can get long, especially when using long chains of pipes ($\rightarrow$ even more time ...)
data.table has short and expressive code, similar to high-level programming languages

Introducing 'data.table'

What is data.table?

1) A selection of functions, optimised to work efficiently with large amounts of data (e.g., fread, fwrite, or for reshaping long-wide / wide-long)

2) A separate dialect for R, with some degree of similarity to SQL or C/C++ code

3) A unique chance to learn about R and computing, with the aim of improving processes and futureproofing our work

Source: https://rdatatable.gitlab.io/data.table/

Introducing 'data.table'

Standard 'data.table' notation

Standard format 1:   data[i , j]   → basic format when not operating on groups

Standard format 2:   data[i , j , by]   → when operating by group

"data" is our dataset
"i"   subset of "data" based on row information ("subset on observations/rows")
"j"   states what to execute for the columns   ("execute on variables/columns")
"by"   defines whether "i" and "j" should be done by groups

Introducing 'data.table'

How to read data.table?

data[sex == "Male",] → subset where sex is "Male", nothing to execute

data[, V2 := V1+1] → nothing to subset, create a new variable "V2" which is "V1+1"

data[sex == "Male", V2 := V1+1] → subset where sex is "Male", then create V2…

data[sex == "Male", .(V2 = max(V1))] → subset where sex is "Male", then return a new variable

data[, .(V2 = max(V1)), by = c("sex")] → no subset, return new variable by group ("sex" variable)

Introducing 'data.table'

Any guesses?

Situation: A three-variable dataset (ID, age, income) reflecting yearly income data (repeated obs.)

data2 <- data[age >= 16, .(income_median = median(income)), by = c("ID")]

Introducing 'data.table'

Any guesses?

Situation: A three-variable dataset (ID, age, income) reflecting yearly income data (repeated obs.)

```
data2 <- data[age >= 16, .(income_median = median(income)), by = c("ID")]
```

… assign to a new object data2 something that comes out of data, new data will have 1 row per ID

… subset is what comes before the comma: select those aged 16 or older

… the new variable "income_median" is the median income of all recorded incomes of the subset

… which was established separately for all "IDs" (e.g., there are multiple income records per "ID")

Hands-on Session

All course materials are available on GitHub!

No account required, a .zip bundle can be downloaded

Repository: 2024_RDS_DT

https://github.com/AndreasxHoehn

Hands-on Session

Objectives: Big Picture + Transferable Skills + Strong Foundation

1) Benchmarking time: 'microbenchmark::microbenchmark()'

2) Benchmarking memory: 'object.size()' and variable types

3) Tracing the location of objects within 'tracemem()'

4) Introduction to 'data.table' - basic functions, subsets, creating new variables, group by operations, reshaping data "wide to long" & "long to wide"

Further Resources

data.table on cran: https://cran.r-project.org/web/packages/data.table/

Benchmarking data.table operations: https://tysonbarrett.com//jekyll/update/2019/10/06/datatable_memory/

Benchmarking joins: https://tysonbarrett.com/jekyll/update/2019/10/11/speed_of_joins/

A good basic data.table intro: https://atrebas.github.io/post/2020-06-17-datatable-introduction/

The official data.table FAQ: https://cran.r-project.org/web/packages/data.table/vignettes/datatable-faq.html

Q&A

Questions – Comments – Feedback?

THANK YOU FOR LISTENING

Find out more

Website: www.sipher.ac.uk

Email: sipher@Glasgow.ac.uk

Twitter: @SipherC